

The NoiseSocket Protocol

Alexey Ermishkin (scratch@virgilsecurity.com)

Revision 0draft, 2017-07-12

Contents

Abstract	1
1. Overview	2
2. Messages	2
2.1. Handshake message	2
2.2. Transport message	2
2.2.1. Plaintext payload and padding	3
2.3. Negotiation data	3
3. Prologue	3
4. Running the protocol	4
5. API	4
6. Examples	5
7. IPR	6
7. References	7

Abstract

NoiseSocket is an extension of the Noise Protocol Framework (developed by the authors of Signal and currently used by WhatsApp) that enables quick and seamless secure link between multiple parties with minimal code space overhead, small keys, and extremely fast speed. It uses raw public keys, modern AEAD ciphers and hash functions. It can be used where X.509 infrastructure isn't required and targets IoT devices, microservices, back-end applications such as datacenter-to-datacenter communications.

1. Overview

NoiseSocket describes the packet structure and processing rules for **handshake** and **transport** stages.

Usually 3 or 2 messages are transmitted during the **handshake** which contain public keys as well as optional data like chosen protocol parameters, certificates, signatures and so on. For the Noise protocol framework to be able to “understand” the message contents, it must know all protocol parameters beforehand. So, every handshake packet contains an optional field called ****negotiation_data**** which is not a part of the handshake message and is transmitted in plaintext. To protect the initial ****negotiation_data**** from modifications we make it a part of the **Noise Prologue**. Usage of ****negotiation_data**** without making it a part of the handshake is insecure.

The **transport** phase is responsible for transmitting the actual data which will be encrypted and optionally padded so that packet contents would not be easily distinguished from each other.

2. Messages

There are two types of messages which differ by structure:

- Section 1.1 Handshake message
- Section 1.2 Transport message

All numbers are unsigned big-endian.

2.1. Handshake message

For the simplicity of processing, all handshake messages have identical structure:

- negotiation_data_len (2 bytes)
- negotiation_data
- noise_message_len (2 bytes)
- noise_message

negotiation_data_len and noise_message_len are 2-byte unsigned integers that store the number of bytes for the corresponding negotiation_data and noise_message fields.

2.2. Transport message

Each transport message has a special ‘data_len’ field inside its plaintext payload, which specifies the size of the actual data. Everything after the data is considered

padding. **Padding contents are arbitrary, and must be ignored by the recipient.**

65517 is the max value for data_len: 65535 (noise_message_len) - 16 (for authentication tag) - 2 (for data_len field itself)

The encrypted packet has the following structure:

- packet_len (2 bytes)
- encrypted data

2.2.1. Plaintext payload and padding

Plaintext payload has the following structure:

- data_len (2 bytes)
- data
- remaining bytes: padding

2.3. Negotiation data

The negotiation_data field is used to identify the protocols used, versions, signs of using a fallback protocol and other data that must be processed before reading the actual noise_message.

Though it can be present in every handshake message, it can be used safely only when it is included into the Noise handshake through Prologue or other mechanisms like calling MixHash() before writing the message.

3. Prologue

Client uses following extra data and fields from the first message to calculate the Noise prologue:

```
-> "NoiseSocketInit" string || negotiation_data_len || negotiation_data
```

where || denotes concatenation

If server decides to start a new protocol instead of responding to the first handshake message, it calculate the Noise prologue using the **full first message contents** plus the length and negotiation_data of its own response. String identifier also changes to “NoiseSocketReInit”.

Thus the prologue structure:

```
<- "NoiseSocketReInit" string || first_handshake_message_contents  
|| negotiation_data_len || negotiation_data
```

where `first_handshake_message_contents` is a concatenation of:

```
negotiation_data_len || negotiation_data || noise_message_len || noise_message
```

4. Running the protocol

They are sent according to the corresponding Noise protocol.

To implement Noise_XX 3 messages need to be sent:

```
-> ClientHello
<- ServerAuth
-> ClientAuth
```

2 for Noise_IK

```
-> ClientHello
<- ServerAuth
```

3 If Fallback is used:

```
-> ClientHello
<- ServerHello
-> ClientAuth
```

5. API

Client and server calls these in sequence.

Initialize:

- INPUT: dh, cipher, hash
- OUTPUT: session object

WriteHandshakeMessage:

- INPUT: [negotiation_data][, cleartext_body]
 - negotiation_data is zero-length if omitted
 - cleartext_body is zero-length if omitted
- OUTPUT: handshake_message

PeekHandshakeMessage:

- INPUT: handshake_message
- OUTPUT: negotiation_data

ReadHandshakeMessage:

- INPUT: handshake_message
- OUTPUT: message_body

After WriteClientAuth / ReadClient, both parties can call Write and Read:

Write:

- INPUT: transport_body[, padded_len]
 - padded_len is zero (no padding) if omitted
- OUTPUT: transport_message

Common logic for the padded_len is follows:

packet_len must be a multiple of the padded_len. To achieve this, you can apply the following steps:

```
maxPacketSize = 65517 /* See 2.2. Transport message*/
```

```
for chunks of data sized maxPacketSize or less do {
```

```
    if padded_len != 0{
```

```
        packet_raw_len = 2 (data_len field) + transport_body_len + 16 (auth tag)
```

```
        needed_padding = padded_len - packet_raw_len % padded_len
```

```
        if (packet_raw_len + needed_padding) > maxPacketSize {
```

```
            needed_padding = maxPacketSize - packet_raw_len /* fill up to the packet size*/  
        }
```

```
        ..add needed_padding bytes to the current chunk
```

```
    }
```

Read:

- INPUT: transport_message
- OUTPUT: transport_body

6. Examples

An example first message negotiation_data which allows to determine which algorithms and pattern were used:

- version_id (2 bytes) (has value 1 by default)
- pattern_id (1 byte)
- dh_id (1 byte)
- cipher_id (1 byte)
- hash_id (1 byte)

pattern_id, dh_id, cipher_id and hash_id can be taken from the Noise-c implementation

For example, NoiseXX_25519_AESGCM_SHA256 would be

- pattern_id : 9
- dh_id : 1
- cipher_id : 2
- hash_id : 3

An example of second message negotiation data:

- version_id (2 bytes) (usually same as client)
- status_id (1 byte) (0 if handshake continues, 1 if fallback, 0xFF if server does not understand client)

7. IPR

The NoiseSocket specification (this document) is hereby placed in the public domain.

7. References